

# lua\_iphtable reference

git.pdh

## **lua\_iphtable.h**

### **Defines**

**LUA\_IPTABLE\_ID**

Identity for the `table_t`-userdata.

**LUA\_IPT\_ITR\_GC**

Identity for the `itr_gc_t`-userdata.

**ipt ERROR numbers**

## **lua\_iphtable.c**

Lua bindings for iphtable

### **Structures**

**itr\_gc\_t**

The `itr_gc_t` type has 1 member: `table_t *t` and serves only to create a userdata to be supplied as an upvalue to each iterator function. That userdata has a metatable with a `__gc` garbage collector, (see `ipt_itr_gc`) which will remove radix nodes in `t` that are (still) flagged for deletion only when it is safe to do so (i.e. `t->itr_lock` has reached zero).

## lualib

### luaopen\_ipatable

```
int luaopen_ipatable(lua_State *L);
```

Called by Lua upon `require("ipatable")` to initialize the module. Returns `ipatable`'s module-table with the module functions and constants. Instance methods go into its metatable.

## error handlers

### lipt\_error

```
static int lipt_error(lua_State *L, int errno);
```

A helper function to set L's global error numer, clear the stack, push nil, and return 1. May be used by functions to bail out while setting an error indicator.

## stack functions

- `iptL`-functions get/set/push stack values.
- `iptT`-functions manipulate k,v-pairs of a table on top of L

### iptL\_gettable

```
static table_t *iptL_gettable(lua_State *L, int index);
```

Checks whether the L stack value at `index` contains a userdata of type `LUA_IPTABLE_ID` and returns a `table_t` pointer. Errors out to Lua if the stack value has the wrong type.

### iptL\_getaf

```
static int iptL_getaf(lua_State *L, int idx, int *af);
```

Checks whether `L[idx]` is a known `af_family` number or name and sets `*af` accordingly or to `AF_UNSPEC` if an unknown `af_family` was found. If `L[idx]` is missing, `*af` is not modified. Caller decides if a missing or unknown `af_family` is an error or not.

### `iptL_getbinkey`

```
static int iptL_getbinkey(lua_State *L, int idx, uint8_t *buf, size_t *len);
```

Copy a binary key, either an address or a mask, at given ‘idx’ into given ‘buf’. ‘buf’ size is assumed to be MAX\_BINKEY. A binary key is a byte array [LEN | key bytes] and LEN is total length. However, radix masks may have set their LEN-byte equal to the nr of non-zero bytes in the array instead.

### `iptL_getpfxstr`

```
const char * iptL_getpfxstr(lua_State *L, int idx, const char **, size_t *len)
```

Checks if L[idx] is a string and yields a const char ptr to it or NULL. Sets len to the string length reported by Lua’s stack manager. Lua owns the memory pointed to by the return value (no free by caller needed).

### `iptL_refpcreate`

```
static int iptL_refpcreate(lua_State *L);
```

Create a reference id for the lua\_State given. This allocates an int pointer, gets a new ref\_id, stores it in the allocated space and returns a pointer to it.

### `iptL_refpdelete`

```
static void iptL_refpdelete(void *L, void **r);
```

Delete a value from LUA\_REGISTRYINDEX indexed by **r (which is treated as an int)**. Function signature is as per `purge_f_t` (see `iptable.h`) and acts as the table’s purge function to release user data.

### `iter_fail_f`

```
static int iter_fail_f(lua_State *L);
```

An iteration func that terminates any iteration.

## module functions

### `iptable.new`

```
static int ipt_new(lua_State *L);
```

Creates a new userdata, sets its `iptable` metatable and returns it to Lua. It also sets the purge function for the table to `iptL_refpdelete` which frees any memory held by the user's data once a prefix is deleted from the radix tree.

#### **`iptable.tobin`**

Return byte array for binary key, masklength & AF for a given prefix, nil on errors. Note: byte array is [LEN | key-bytes], where LEN is the total length of the byte array.

#### **`iptable.tostr`**

Return string representation for a binary key, nil on errors

#### **`iptable.masklen`**

Return the number of consecutive msb 1-bits, nil on errors. *note: stops at the first zero bit, without checking remaining bits.*

#### **`iptable.size`**

Returns the number of hosts in a given prefix, nil on errors

- uses Lua's arithmetic ( $2^{\text{hostbits}}$ ), since ipv6 can get large.

#### **`iptable.address`**

Return host address, masklen and af\_family for pfx; nil on errors.

```
addr, mlen, af = iptable.address("10.10.10.10/24")
```

#### **`iptable.network`**

Return network address, masklen and af\_family for pfx; nil on errors.

#### **`iptable.neighbor`**

Given a prefix, return the neighbor prefix, masklen and af\_family, such that both can be combined into a supernet with masklen -1; nil on errors. Note that a /0 will not have a neighbor prefix with which it could be combined.

### **iptable.incr**

Return address, masklen and af\_family by adding an offset to a pfx; If no offset is given, increments the key by 1. Returns nil on errors. Note: any mask is ignored with regards to adding the offset.

### **iptable.decr**

Return address, masklen and af\_family by adding an offset to a pfx; If no offset is given, increments the key by 1. Returns nil on errors. Note: any mask is ignored with regards to adding the offset.

### **iptable.invert**

Return inverted address, masklen and af\_family. Returns nil on errors. Note: the mask is NOT applied before inverting the address.

### **iptable.reverse**

Return a reversed address, masklen and af\_family. Returns nil on errors. Note: the mask is NOT applied before reversing the address.

### **iptable.broadcast**

Return broadcast address, masklen and af\_family for pfx; nil on errors.

### **iptable.mask**

```
mask, af, mlen = iptable.mask(af, mlen, invert?)  
  
static int ipt_mask(lua_State *L); // <-- [af, mlen, invert?]
```

Returns the mask, as a string, for the given address family **af** and mask length **mlen**. Inverts the mask if the optional argument **invert** evaluates to true.

### **iptable.hosts**

Iterate across host addresses in prefix pfx.

The optional **incl** argument defaults to false, when given & true, the network and broadcast address will be included in the iteration results. If pfx has no mask, the af's max mask is used.

### `iter_hosts_f`

The `iter_f` for `iptable.hosts(pfx)`, yields the next host ip until stop value is reached. Ignores the stack: it uses upvalues for next, stop

### `iptable.interval`

```
static int ipt_interval(lua_State *L) // <-- [start, stop]
for pfx in iptable.interval("10.10.10.0", "10.10.10.8") do
    print(pfx)
end
-- would print
10.10.10.0/29
10.10.10.8/32
```

Iterate across the prefixes that, combined, cover the address space between the start & stop addresses given (inclusive). Any masks in either `start` or `stop` are ignored and both need to belong to the same AF\_family.

## instance methods

### `iptm_gc`

Garbage collect the table: free all resources Assigned to metatable's `__gc` method.

### `iptm_newindex`

Implements `t[k] = v`

If `v` is nil, `t[k]` is deleted. Otherwise, `v` is stored in the `lua_registry` and its ref-value is stored in the radix tree. If `k` is not a valid ipv4/6 prefix, cleanup and ignore the request.

### `iptm_index`

Given an index `k`:

- do longest prefix search if `k` is a prefix without mask,
- do an exact match if `k` is a prefix with a mask
- otherwise, do metatable lookup for property named by `k`.

#### **iptm\_len**

Return the sum of the number of entries in the ipv4 and ipv6 radix trees as the 'length' of the table.

#### **iptm\_tostring**

Return a string representation of the iptable instance (with iptm\_counts).

#### **iptm\_counts**

Return both ipv4 and ipv6 iptm\_counts as two numbers.

#### **iter\_kv**

Setup iteration across key,value-pairs in the table. The first ipv4 or ipv6 leaf node is pushed. The iter\_pairs iterator uses nextleaf to go through all leafs of the tree. It will traverse ipv6 as well if we started with the ipv4 tree first.

#### **iter\_kv\_f**

The iter\_f for \_\_iter\_kv(), yields all k,v pairs in the tree(s).

- upvalue(1) is the leaf to process.

#### **lipt\_itermore**

Iterate across more specific prefixes. AF\_family is deduced from the prefix given. Note this traverses the entire tree if pfx has a /0 mask.

#### **lipt\_itermoref**

Iterate across more specific prefixes one at a time.

- top points to subtree where possible matches are located
- rn is the current leaf under consideration

#### **t:less**

Iterate across prefixes in the tree that are less specific than pfx.

- search may include pfx itself in the results, if `incl` is true.
- `iter_less_f` takes (pfx, mlen, af)

#### **iter\_less\_f**

Iterate across less specific prefixes relative to a given prefix.

#### **t:masks**

Iterate across all masks used in AF\_family's radix tree. Notes:

- a mask tree stores masks in `rn_key` fields.
- a mask's `KEYLEN` == nr of non-zero mask bytes, rather `LEN` of byte array.
- a /0 (zeromask) is never stored in the radix mask tree (!).
- a /0 (zeromask) is stored in `ROOT(0.0.0.0)`'s *last* dupedkey-chain leaf.
- the AF\_family cannot be deduced from the mask in `rn_key`.

So, the `iter_f` needs the AF to properly format masks for the given family (an ipv6/24 mask would otherwise come out as 255.255.255.0) and an explicit flag needs to be passed to `iter_f` that a zeromask entry is present.

#### **iter\_masks\_f**

`iter_f` for `t:masks(af)`, yields all masks used in af's radix mask tree. Masks are read-only: the Lua bindings donot (need to) interact directly with a radix mask tree.

#### **t:merge**

Iterate across groups of pfx's that may be combined, to setup just return the first leaf of the correct AF\_family's tree.

#### **iter\_merge**

Iterate across prefixes that may be combined.

Returns the supernet and a regular table with 2 or 3 pfx,value entries.

Caller may decide to keep only the supernet or parts of the group. If the supernet exists in the tree it will be present in the table as well as a third entry.



The next node stored as upvalue before returning, must be the first node of the next group.

#### **iptL\_pushitr\_gc**

During iteration(s), delete operations only flag radix nodes (leafs holding a prefix) for deletion so they are not actively removed, only flagged for deletion. This way, the ‘next leaf’ node, stored as an upvalue by the iter\_XXX functions, is guaranteed to still exist on the next iteration.

Such ‘inactive’ leafs are still part of the tree, but are seen as absent by regular tree operations.

The iterator garbage collector only actually removes those flagged for deletion when there are no more iterators active (ie. itr\_count reaches zero).

#### **ipt\_itr\_gc**

The garbage collector for iterator functions, pushed by the factory function and executed when Lua decides to do some garbage collection.

#### **t:radix**

Iterate across all nodes of all types in af’s radix trees.

Only includes the radix\_mask tree if the optional maskp argument is true; The iterator will return the C-structs as Lua tables.

#### **t:lipt\_iterradixf**

- return current node as a Lua table or nil to stop iteration
- save next (type, node) as upvalues (1) resp. (2)

#### **iptT\_setfstr**

TODO: doc and move

#### **iptT\_setint**

TODO: doc and move

**iptT\_setkv**

set (rn->rn\_key, v) on Table on top of L, based on pfx provided and its value stored (indirectly) in the corresponding radix tree. TODO: doc, move

**iptL\_pushrn**

push\_rdx\_'s each translating a struct to a lua table. Will translate and push IPTF\_DELETE'd nodes for completeness sake. The lua caller may decide herself what to do with those.

**iptL\_pushrn**

TODO: doc, move

**iptL\_pushrh**

TODO: doc, move

push a radix head onto L (as a table)

**iptL\_pushrmh**

TODO: doc, move push a radix mask head onto L (as a table)

**iptL\_pushrm**

TODO: doc, move